

# Retrieval-Augmented Generation: Architecture, Evaluation, and Future Directions

**Ava (Zheer) Wang**

Georgia Institute of Technology  
zwang3657@gatech.edu

**Minghao Zheng**

Georgia Institute of Technology  
minghao.zheng@gatech.edu

**Ruishu Cao**

Georgia Institute of Technology  
rcao70@gatech.edu

**Devansh Khunteta**

Georgia Institute of Technology  
dkhunteta3@gatech.edu

## Abstract

Retrieval-Augmented Generation (RAG) grounds large language models (LLMs) in external knowledge, yet current practice remains fragmented across system architectures, evaluation protocols, and deployment settings. This paper presents a unified, practice-oriented analysis of RAG pipelines that connects design choices in preprocessing, retrieval, and generation to downstream faithfulness and efficiency. We further position RAG relative to fine-tuning, clarifying when external retrieval or internal adaptation is most effective for knowledge-intensive tasks. Drawing on deployments in healthcare, finance, legal, and enterprise systems, we identify persistent challenges, such as retrieval drift, latency–cost trade-offs, and privacy, and discuss emerging future directions including multi-agent RAG and tool-integrated retrieval. Our goal is to provide a coherent framework for building reliable, transparent, and continuously updatable knowledge-aware systems.

## 1 Introduction

Large language models (LLMs) have achieved strong performance across NLP tasks, but their knowledge remains static and opaque: they cannot incorporate newly emerging facts or organization-specific information without retraining, and they may generate confident yet unsupported statements when relevant context is missing (Ji et al., 2023). These limitations hinder deployment in domains where information changes rapidly and provenance is essential.

Retrieval-Augmented Generation (RAG) mitigates these issues by conditioning generation on evidence retrieved from an external, updatable knowledge base (Lewis et al., 2020). By decoupling knowledge storage from parametric reasoning, RAG improves factuality, transparency, and adaptability. Despite its rapid adoption, practical development of RAG systems remains fragmented:

design choices in preprocessing, chunking, embedding, retrieval, reranking, and prompting vary widely across implementations, and evaluation protocols provide inconsistent signals about faithfulness, relevance, and efficiency. Real-world reports further highlight persistent challenges such as retrieval drift, latency–cost trade-offs, and privacy concerns.

These gaps motivate the need for a unified, practice-oriented perspective on RAG system design and evaluation. In this paper, we structure the analysis around three key dimensions:

- (i) **Architecture and design:** a systematic overview of preprocessing, retrieval, and generation components in RAG system
- (ii) **Evaluation and optimization:** an integrated view of retrieval quality, grounding faithfulness, and end-to-end system efficiency, together with the optimization levers these metrics suggest.
- (iii) **Applications and lessons:** a synthesis of deployed systems across conversational assistants, healthcare, finance, legal, and enterprise settings, including common failure modes and emerging multi-agent workflows.

Overall, our analysis clarifies how RAG transforms static LLMs into verifiable and continuously updatable *knowledge-aware* systems, and surfaces open challenges that must be addressed for reliable deployment at scale.

## 2 Background and Motivation

LLMs encode world knowledge in their model parameters, enabling strong generalization but imposing fundamental limitations: their internal “memory” is static, costly to update, and incapable of integrating post-training information. Consequently, pretrained LLMs struggle to incorporate

organization-specific content, dynamic regulations, or temporally sensitive facts without retraining, a process that is computationally expensive and operationally inflexible, thereby limiting practical deployment.

RAG mitigates these limitations by complementing parametric memory with non-parametric access to external knowledge sources. At inference time, the model retrieves relevant documents from a continuously updatable index, rather than relying solely on its fixed internal representations (Lewis et al., 2020). This retrieval step grounds generation in verifiable evidence, improving factual consistency and reducing hallucinations (Ji et al., 2023). Crucially, because the knowledge base is externalized, system behavior can be updated by modifying the indexed corpus or embeddings—without retraining the underlying language model.

RAG has demonstrated strong performance in domain-specific settings, where accuracy, regulatory compliance, and source traceability are critical (Sharma et al., 2024). However, practical deployments highlight that RAG systems involve a range of design decisions—such as how to split documents into chunks, which embedding models to use, how many results to retrieve, how to re-rank them, and how to construct prompts. Each of these choices can significantly affect the factual reliability and efficiency of the system.

These observations suggest that RAG is not only a conceptual approach for accessing external knowledge, but also a practical system architecture that requires thoughtful design and systematic evaluation. This motivates our practice-oriented analysis of architectural trade-offs, evaluation strategies, and emerging challenges.

## 3 RAG Architecture

### 3.1 Preprocessing (Offline)

In the preprocessing phase, source documents are segmented into smaller, semantically coherent units to enable fine-grained retrieval. Each chunk is encoded using a pretrained embedding model, producing dense vector representations that preserve semantic similarity. These embeddings are indexed in a vector database (e.g., FAISS, Pinecone, Weaviate), which supports efficient approximate nearest neighbor search using similarity metrics such as cosine distance or inner product. Metadata such as document identifiers, chunk indices, and source information is typically stored alongside the em-


beddings to facilitate traceability and post-retrieval filtering.

### 3.2 Retriever

The retriever component is responsible for identifying semantically relevant document chunks from the vector index in response to a user query. At inference time, the query is encoded into a dense vector using the same embedding model used during preprocessing, ensuring alignment between the query and stored representations. The system then computes similarity scores between the query vector and all stored embeddings, and selects the top- $k$  chunks with the highest scores as candidate contexts. These retrieved chunks may be further filtered or re-ranked based on additional criteria such as relevance heuristics or source reliability, prior to being passed to the generator.

### 3.3 Generator (LLM)

In the final stage, the system constructs a prompt by concatenating the retrieved context with the user query. This prompt is then passed to a pretrained large language model (e.g., GPT-4, LLaMA) to produce a contextually grounded response. By conditioning generation on retrieved content, the model is able to incorporate factual information that lies outside its training distribution, thereby improving the accuracy and trustworthiness of its outputs while mitigating hallucinations.



**Figure 1: RAG workflow.** Raw documents are first chunked and converted into dense vector embeddings, which are stored in a vector database. At inference time, a user query is embedded using the same model and used to retrieve the most relevant chunks based on semantic similarity. These retrieved chunks are combined with the query to form a prompt, which is then processed by a large language model (LLM) to generate a grounded, context-aware answer.

*Image source: Daily Dose of Data Science*

## 4 RAG System Evaluation Metrics

Evaluating RAG systems requires a multi-dimensional perspective, as system behavior is shaped not only by the quality of generation but also by the effectiveness of retrieval and the efficiency of the overall serving pipeline (Yu et al., 2025). In this work, we evaluate RAG systems

along three key dimensions: retrieval quality, generation quality, and end-to-end system efficiency.

#### 4.1 Retrieval Quality

Evaluating retrieval quality involves assessing whether the system retrieves the most relevant information needed to generate an accurate and grounded response. Commonly used metrics include Recall@k, Precision@k, Mean Reciprocal Rank (MRR), and Normalized Discounted Cumulative Gain (nDCG) (Xiong et al., 2020), which capture both how much relevant content is retrieved and how well it is ranked. High recall is especially important, as missing key evidence at this stage limits the generator’s ability to produce faithful answers. To better understand how retrieved content contributes to generation, recent work has proposed metrics such as the Context Utilization Rate (Hagström et al., 2025), which measures the proportion of retrieved text actually used in the final response. This provides a more detailed view of the interaction between the retriever and the generator.

#### 4.2 Generation Quality

The generation stage is typically evaluated based on three key aspects: fluency, semantic fidelity, and factual faithfulness. Traditional reference-based metrics such as BLEU (Papineni et al., 2002), ROUGE-L (Lin, 2004), and BERTScore (Zhang et al., 2020) assess lexical or embedding-level similarity to reference outputs. However, these metrics are limited in their ability to detect hallucinations or unsupported factual claims. To address this gap, recent work has introduced faithfulness-centered evaluation methods, including QAFactEval (Fabri et al., 2022), FactScore (Min et al., 2023), and LLM-as-a-Judge frameworks (Zheng et al., 2023). These approaches compare generated content directly against retrieved evidence, yielding interpretable scores for factual precision and hallucination rate. Such metrics are especially critical in knowledge-intensive domains, where factual accuracy is prioritized over surface-level similarity.

#### 4.3 End-to-End Deployment Quality

Model and human based holistic evaluations, such as those using LLM as a Judge scoring, integrate multiple dimensions including relevance, coherence, and faithfulness into a unified quality score. These evaluations provide a comprehensive understanding of RAG system behavior that isolated metrics often fail to capture.

Beyond accuracy, real-world RAG deployment also depends on system level efficiency. End to end latency, throughput, and GPU cost per query are widely used to assess computational overhead and scalability (Kim and Mahajan, 2025). The ratio between retrieval and generation time reflects the balance of pipeline components, while GPU utilization and energy consumption quantify cost efficiency.

These metrics are inherently interdependent. Improving Recall@k generally increases faithfulness but may lead to higher latency, while aggressive token pruning can reduce cost but risks omitting essential evidence. As a result, evaluation metrics not only measure performance but also implicitly define the trade offs and optimization goals of a RAG system. A well designed evaluation framework is essential for guiding effective optimization strategies.

#### 4.4 Optimization in RAG System

Guided by metric feedback, RAG optimization aims to jointly maximize faithfulness and efficiency. Recent literature has converged on four complementary directions that target different stages of the RAG pipeline.

First, self-knowledge triggering focuses on when to invoke retrieval. Confidence estimators determine whether external augmentation is necessary based on model uncertainty. By avoiding retrieval in high-confidence scenarios, systems reduce latency and GPU cost, while selectively retrieving under uncertainty enhances answer accuracy and faithfulness (Zubkova et al., 2025).

Second, optimization explores where to augment within the model architecture. Knowledge can be integrated at different stages: input-level augmentation improves factual grounding by enriching the prompt with relevant evidence (Ram et al., 2023); intermediate-level approaches, such as Retro, enhance contextual reasoning by inserting information into hidden states (Borgeaud et al., 2022); and output-level methods like post-verification help ensure answer consistency and factual accuracy (Dhuliawala et al., 2024).

Third, multi-hop and iterative augmentation aim to improve Recall@k and reasoning completeness for complex queries (Trivedi et al., 2020). These approaches expand the retrieval horizon by chaining multiple queries or steps, but introduce additional latency. Thus, optimization must carefully balance completeness against computational cost.

Finally, cost-efficiency techniques address the growing demand for scalable deployment. Methods such as token pruning (Ratner et al., 2023), context compression (Verma, 2024), and KV-cache reuse (Yang et al., 2025) have been proposed to reduce memory and compute requirements without degrading output quality. These improvements directly enhance throughput and energy efficiency, contributing to more sustainable and scalable RAG systems.

Metrics define the quantitative compass of system behavior, while optimization strategies translate those signals into algorithmic improvements. Through this iterative cycle, modern RAG research advances toward systems that are not only more faithful and accurate but also computationally efficient and adaptive.

## 5 Real-World Applications

RAG has evolved from a research concept into a foundational paradigm underpinning many deployed systems. By integrating the reasoning capacity of large language models (LLMs) with dynamic, verifiable retrieval, RAG enables generation over live, domain-specific, and evidence-grounded information. This section surveys representative use cases across conversational assistants, healthcare, enterprise analytics, and emerging multi-agent decision-support environments.

### 5.1 Conversational and Search Assistants

RAG forms the backbone of several modern conversational systems, including web-scale assistants such as Bing Copilot and Perplexity. These agents pair dense retrieval with generative reasoning to access continuously updated knowledge while maintaining coherent dialogue. Unlike static LLMs, RAG-based assistants deliver referenceable and citation-supported responses in real time (Lewis et al., 2020; Gao et al., 2023). Beyond open-domain chat, similar architectures have been applied to customer-support automation, internal knowledge-base querying, and educational tutoring, settings where content evolves quickly and factual reliability is critical.

### 5.2 Healthcare and Biomedical Applications

In clinical and biomedical domains, RAG supports factual grounding and interpretability, two properties essential for responsible AI adoption. Recent work demonstrates that RAG pipelines

can generate patient-specific summaries, assist in trial-report drafting, and deliver decision support anchored in medical literature. For example, a dementia-care RAG system combined with parameter-efficient fine-tuning (LoRA/QLoRA) produced context-aware guidance for caregivers while markedly reducing hallucinations (Al-Qudah et al., 2025b). Similarly, systems integrating electronic health records and PubMed retrieval achieved verifiable clinical-trial analyses and reduced report-writing time by over 70% (Al-Qudah et al., 2025a; Wang et al., 2024). These results highlight RAG’s capacity to combine retrieval grounding with lightweight adaptation for accurate, cost-effective, and explainable medical NLP.

### 5.3 Domain-Specific and Enterprise Systems

RAG has also gained traction in enterprise contexts demanding factual precision and regulatory compliance. In finance, domain-adapted pipelines, such as those modeled on FinanceBench, retrieve regulatory filings, earnings reports, and market data to produce numerically consistent analyses. Legal and scientific assistants employ similar retrieval-generation hybrids to synthesize case law and research findings with transparent citation and attribution (Sharma et al., 2024; Patel et al., 2024). Within organizations, RAG facilitates secure querying of proprietary documents without retraining, supporting compliance audits, policy updates, and technical troubleshooting based on continuously refreshed internal knowledge.

### 5.4 Multi-Agent Decision-Support Systems

Recent advances extend RAG toward multi-agent frameworks where specialized agents, retrievers, planners, synthesizers, and verifiers, collaborate to iteratively refine both evidence and reasoning. Deployments such as university-admissions counseling systems demonstrate that decomposing roles across agents can reduce hallucination by over 90% and lower response latency while preserving interpretability (Nguyen-Duc et al., 2025). This transition from monolithic retrieval-generation pipelines to modular, cooperative reasoning architectures marks a step toward scalable and auditable knowledge reasoning in complex decision-support environments.

## 6 RAG vs. Fine-Tuning

While both RAG and fine-tuning extend large language models (LLMs) with domain-specific knowl-

| Aspect             | RAG                            | Fine-Tuning               |
|--------------------|--------------------------------|---------------------------|
| Mechanism          | Retrieve external knowledge    | Update model weights      |
| Knowledge Type     | Non-parametric (outside model) | Parametric (inside model) |
| Update Cycle       | Real-time refresh of index     | Requires retraining       |
| Transparency       | Provides citations / evidence  | Opaque, black-box memory  |
| Cost & Scalability | Low; easily updated            | High; resource-intensive  |

Table 1: RAG vs. Fine-Tuning — Two Paths to Model Adaptation

edge, they differ fundamentally in how information is incorporated, updated, and validated. Recent empirical comparisons across industrial and academic settings clarify their complementary roles and trade-offs (Lewis et al., 2020; Gao et al., 2023; Han et al., 2023; Kumar et al., 2024b).

### 6.1 Knowledge Integration Mechanism

Fine-tuning embeds new information directly into a model’s parameters, enabling implicit recall of linguistic and factual patterns. However, this parametric memory is static, once training is complete, the model cannot reflect newly emerging data without further fine-tuning. RAG instead employs a non-parametric memory through external retrieval: rather than encoding facts in weights, it dynamically conditions generation on documents drawn from an updatable index (Lewis et al., 2020). Empirical evidence supports these distinctions: in Microsoft’s agriculture study, fine-tuning improved question-answer accuracy by roughly six percentage points, while augmenting the same model with RAG produced an additional five-point gain, showing that retrieval complements internal adaptation rather than replacing it (Kumar et al., 2024b).

### 6.2 Adaptability and Cost Efficiency

Fine-tuning demands curated data, GPU compute, and per-domain checkpoints, making maintenance costly for organizations managing multiple or rapidly changing domains. RAG decouples knowledge from parameters: refreshing the document store or re-embedding new material immediately updates model behavior without retraining. This design yields superior scalability for dynamic, data-rich environments such as finance, healthcare, and enterprise analytics (Sharma et al., 2024; Al-Qudah et al., 2025a). A comparative chatbot

evaluation by Kumar et al. (2024a) confirmed this distinction—RAG (implemented with LangChain and LLaMA-3) achieved deployment within hours and matched a fine-tuned DistilBERT baseline in semantic accuracy while consuming markedly less compute.

### 6.3 Faithfulness and Interpretability

Fine-tuned models generate from internal representations and provide no direct trace to supporting evidence, which complicates verification. RAG, by contrast, grounds its outputs in retrieved documents, enabling explicit citation and auditability (Gao et al., 2023). Both Microsoft’s and Kumar et al.’s experiments observed lower hallucination rates and higher factual alignment for RAG pipelines, though retrieval drift, returning semantically similar but irrelevant passages, remains a key source of error (Kumar et al., 2024b,a).

### 6.4 Complementary Use and Emerging Hybrids

Recent work emphasizes that RAG and fine-tuning are complementary rather than competing paradigms. Fine-tuning internalizes stable procedural or stylistic expertise, while RAG provides dynamic factual grounding. Hybrid systems, such as LoRA- or QLoRA-based models augmented with retrieval, achieve the best of both worlds. In the Microsoft pipeline, combining LoRA fine-tuning with retrieval improved geographically localized reasoning, raising answer similarity from 47 % to 72 % (Kumar et al., 2024b). Similarly, domain-specific chatbots combining fine-tuned transformers with RAG produced higher BERTScore precision and Flesch–Kincaid readability than either method alone (Kumar et al., 2024a). Together, these findings reinforce a unifying view: fine-tuning internalizes stable expertise, whereas RAG externalizes and continually refreshes it—forming a cohesive framework for scalable, verifiable, and knowledge-adaptive generation.

## 7 Limitations

Although RAG architectures substantially improve factual grounding, they also introduce several inherent limitations. A central issue is the system’s dependence on retriever quality: the generation is constrained by the documents surfaced at retrieval time. If the retriever yields semantically similar but topically irrelevant passages, the resulting answer reflects a “garbage in, garbage out” failure mode.

Dense embeddings are especially prone to this mismatch, high cosine similarity does not guarantee true relevance, allowing off-topic or low-quality evidence to shape downstream reasoning. Furthermore, if the underlying corpus contains outdated or biased material, those flaws propagate into the model’s outputs, amplifying misinformation risks (Ouyang et al., 2025).

RAG systems also impose architectural and operational complexity. Incorporating an external retrieval call adds latency, as every query must pass through embedding computation and vector-index lookup before generation begins. This overhead complicates interactive, real-time applications and introduces new systems-engineering burdens. Maintaining a large vector database requires continuous ingestion and re-embedding of updated content, with substantial storage and compute costs. Modern similarity-search infrastructure such as FAISS (Johnson et al., 2019) enables scalable indexing, but keeping embeddings fresh and ensuring high recall at scale remains non-trivial.

Despite grounding in retrieved documents, RAG models still hallucinate. When relevant evidence is absent or incomplete, the generator may fill gaps using parametric memory, producing statements that appear justified but lack explicit support. Even when citations are provided, attribution is imperfect: models frequently attach plausible sources that do not verify the specific claims being made. Prior work shows that hallucinations can persist even when retrieval is correct, since the generator may loosely paraphrase or misattribute content (Ji et al., 2023). This creates a false sense of reliability; the presence of retrieved text can mask the fact that parts of the answer remain fabricated.

Privacy and personalization further complicate deployment. Applying RAG over proprietary or user-specific data introduces risks, since indexing sensitive material expands the system’s attack surface. Private facts encoded in the vector store may leak through generation, and current RAG pipelines lack built-in mechanisms for isolating user data or enforcing deletion guarantees (Carlini et al., 2021). While personalization promises tailored retrieval augmented with user context, it remains challenging to incorporate such information safely. Prior work shows that retrieval models can inadvertently expose memorized content or recombine private details unless strong safeguards are applied.

## 8 Challenges and Open Problems

Before exploring specific failure modes in RAG systems, it is useful to contextualize where the community perceives the biggest production bottlenecks. Figure 1 shows results from the 2024 *LangChain State of AI Agents* report, highlighting that performance quality remains the dominant limitation (reported by 41% of respondents), followed by cost, safety, and latency concerns (LangChain, 2024). This trend underscores that improving retrieval relevance and grounding quality is an urgent research priority for real-world RAG deployment.



Figure 1: Reported limitations in deploying multi-agent and RAG systems at scale. Data from the *LangChain State of AI Agents* (2024) survey.

### 8.1 Retrieval Drift and Ranking Bias

A pervasive challenge is retrieval drift, where the system retrieves passages that share lexical or semantic features with the query but do not align with the user’s actual intent. As a result, the model may produce an answer that is technically grounded yet contextually irrelevant, giving an illusion of correctness while missing the true target. Dense indexes are particularly susceptible: near-duplicate or loosely related documents often dominate the top-k results, crowding out the genuinely relevant content. Ranking mechanisms also introduce bias—longer documents, more keyword-dense pages, or well-linked sources may be systematically prioritized regardless of task relevance (Kim et al., 2024). This bias can cause the model to repeatedly draw from a skewed subset of the corpus, reinforcing retrieval drift.

Mitigation strategies include two-stage retrieval with learned re-ranking models (Nogueira and Cho, 2019), diversification techniques such as maximal marginal relevance (Carbonell and Goldstein, 1998), and dynamic query reformulation via agen-

tic or self-reflective prompting (Gao et al., 2022; Yao et al., 2023). However, robustly correcting drift remains unresolved: it requires feedback loops that continuously diagnose retrieval errors and refine search queries in real time.

## 8.2 Latency and Scalability

RAG’s dual process (retrieve-then-generate) inherently incurs extra latency compared to a standard LLM response. For each user query, the system must perform embedding computations and database lookups before even beginning to generate text. This search time can be substantial, especially against a large corpus or when using complex ranking and filtering (Khattab and Zaharia, 2020). Furthermore, the generation step itself becomes slower as more context is fed into the model; long retrieved documents consume the model’s context window and computation, increasing response delay. Thus, there is a fundamental trade-off between answer quality and speed; incorporating more knowledge makes answers better but slower.

Scalability is another challenge. As the knowledge corpus grows to millions of documents or more, naive retrieval approaches may not meet real-time requirements. Indexing and memory footprint grow with the corpus, and querying a massive index can bottleneck. Techniques like approximate nearest-neighbor search and sharding of the vector index are used to keep retrieval time manageable, but they may sacrifice some recall when tuning for speed (Johnson et al., 2019). For applications serving many users or handling streaming data, maintaining low latency while continuously ingesting new information is an open problem. Potential solutions involve caching frequent queries, using smaller specialized indexes for different topics, or predicting when to skip retrieval for easy questions. Research is ongoing into adaptive retrieval strategies that decide how many documents to retrieve, or whether to retrieve at all, based on query complexity and required precision. Achieving a scalable RAG that feels instantaneous to users is still a frontier in system optimization.

## 8.3 Faithfulness and Attribution

A core promise of RAG is that responses will be grounded in the retrieved evidence. However, ensuring that the generation faithfully uses that evidence is far from solved. RAG models can and do hallucinate, inserting details or assertions not supported by any retrieved document, especially if

the retrieval was imperfect or the query requires synthesis (Ji et al., 2023). Even when the model includes citations or quotes, the correspondence between those references and the claims made can be shaky. Studies highlight that citation correctness is not the only concern; faithfulness is often lacking, meaning the model may produce an answer from parametric memory and then attach a plausible-looking source. In practice, an attributed answer might cite a document that covers the general topic but not the specific fact in the answer.

This issue is closely tied to hallucination in generation. RAG reduces hallucinations compared to unguided LLMs but does not eliminate them; if the retrieval misses relevant detail, the model may fill the gap with guesses. It may also combine pieces of information from different documents in a way that no single source asserts, which makes verification difficult (Ji et al., 2023). Addressing faithfulness will likely require stronger guidance from retrieved text, sentence-level verification or post-hoc checking, and prompting strategies that force explicit source-claim alignment (Asai et al., 2023).

## 8.4 Privacy and Personalization

As RAG systems move from public data toward personal or enterprise applications, privacy becomes a central concern. An open problem is how to give users personalized, context-aware answers without compromising sensitive data. A naive approach to personalization might index all of a user’s documents or chat history into the retriever. While this would allow highly tailored answers, it introduces risks: private information becomes part of the model’s search space and may be exposed inadvertently. Even with access control, the model can leak private details in its output because it treats them as general knowledge (Carlini et al., 2021).

Balancing user-adaptive retrieval with data privacy means the system must incorporate personal data in a secure, consented, and isolated way. One challenge is data isolation in multi-user settings; the retriever should not cross-contaminate data between users. Each user’s context might need a separate index or strict per-user permissions, which adds complexity. Compliance with privacy regulations also matters, including deletion propagation so that removed documents no longer appear in future answers. Future RAG systems will likely employ on-device or federated retrieval, where personal data never leaves a secure device or enclave,

and only non-sensitive embeddings or results are shared (Mao et al., 2025). Differential privacy and encryption can limit leakage when aggregated signals are used to improve ranking.

## 9 Future Directions

RAG research is increasingly moving beyond static, one-shot retrieval toward systems that are adaptive, interpretable, and tightly integrated with external tools. We highlight several emerging directions likely to shape the next generation of knowledge-augmented LLMs.

### 9.1 Structured and Auditable Retrieval via MCP

The Model Context Protocol (MCP) offers a promising foundation for standardized and auditable retrieval. By exposing retrieval as typed, tool-mediated actions rather than an implicit internal step, MCP enables clearer provenance, debuggability, and integration with diverse data sources. This direction aligns with growing interest in tool-augmented LLMs (Schick et al., 2023) and suggests that future RAG systems may rely on unified APIs for querying, filtering, and validating external knowledge, which in turn supports more interpretable and reproducible retrieval traces. Recent work on RAG-MCP (Gan and Sun, 2025) demonstrates the practical value of this approach. By treating tool selection itself as a retrieval problem, RAG-MCP achieves 43.13 percent accuracy compared to 18.20 percent for direct matching and 13.62 percent for a blank baseline, while also reducing average prompt length by over one-third. These results highlight how MCP combined with retrieval can substantially improve tool choice and efficiency in multi-tool environments.

### 9.2 Self-Improving and Continual RAG

Current RAG systems do not leverage the rich feedback signals produced during real use. A key direction is *self-improving* RAG, where user corrections, follow-up queries, and implicit satisfaction signals guide online refinement of both retriever ranking and generation behavior (Ouyang et al., 2022). Iterative retrieval-and-verification loops (Asai et al., 2023; Yao et al., 2023) already reduce hallucination by keeping evidence checking “in the loop.” Extending these ideas to continual learning—while maintaining stability and respecting privacy constraints—remains an open problem, especially in settings involving personal or enterprise data.

### 9.3 Adaptive Retrieval and Dynamic Indexing

Retrieval in most RAG systems remains fixed across queries, even though difficulty and uncertainty vary widely. Adaptive retrieval aims to adjust similarity thresholds, depth, and document budgets based on model confidence or early reasoning signals. Techniques such as Hypothetical Document Embedding (HyDE) (Gao et al., 2022) illustrate how pseudo-answer generation can better align semantic search with user intent. HyDE can significantly improve accuracy by outperforming traditional sparse and dense retrieval methods, especially in zero-shot settings where no fine-tuning is possible. Complementarily, dynamic index maintenance, continuous deduplication, freshness scoring, and re-embedding is essential for preventing drift and staleness in large, evolving corpora. These challenges straddle both retrieval algorithms and systems, raising questions about cost-aware refresh policies and scalable update pipelines.

**Outlook.** Collectively, these directions point toward RAG architectures that are adaptive, verifiable, and increasingly tool-integrated. Rather than relying on static retrieval pipelines, future systems will leverage refreshed retrieval infrastructure to deliver more reliable and interpretable knowledge grounding.

## 10 Conclusion

RAG redefines the boundary between retrieval and generation by turning knowledge access into an adaptive reasoning process. Through our unified analysis of architecture, evaluation, and application, we highlight how faithfulness and efficiency metrics can serve as optimization levers for building reliable knowledge-aware AI. Moving forward, advancing RAG requires integrating multi-agent collaboration, standardized tool protocols, and privacy-preserving retrieval. These directions will push conversational AI from passive recall toward verifiable, self-improving reasoning systems that align accuracy with accountability.

## References

- Omar Al-Qudah, Xin Huang, Li Fang, and 1 others. 2025a. [Automated clinical trial data analysis and report generation by integrating retrieval-augmented generation \(rag\) and large language models](#). *AI*, 6(8).
- Omar Al-Qudah, Rui Zhang, Yifan Liu, and 1 others. 2025b. [Rag pipeline for domain-specific ap-](#)

- lications: A case study in disseminating dementia practices. In *2025 IEEE/ACM International Conference on Connected Health: Applications, Systems and Engineering Technologies (CHASE)*.
- Akari Asai, Jaewoo Jung, Eunsol Choi, and Hannaneh Hajishirzi. 2023. [Self-RAG: Learning to retrieve, generate, and critique for knowledge-intensive NLP](#). *arXiv preprint arXiv:2310.11511*.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, and 9 others. 2022. [Improving language models by retrieving from trillions of tokens](#). *Preprint*, arXiv:2112.04426.
- Jaime Carbonell and Jade Goldstein. 1998. [The use of MMR, diversity-based reranking for reordering documents and producing summaries](#). In *SIGIR*.
- Nicholas Carlini, Florian Tramer, Erik Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, Alina Oprea, and 1 others. 2021. [Extracting training data from large language models](#). In *USENIX Security Symposium*.
- Shehzaad Zuzar Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, and Jason E Weston. 2024. [Chain-of-verification reduces hallucination in large language models](#).
- Alexander Fabbri, Chien-Sheng Wu, Wenhao Liu, and Caiming Xiong. 2022. [QAFactEval: Improved QA-based factual consistency evaluation for summarization](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2587–2601, Seattle, United States. Association for Computational Linguistics.
- Tiantian Gan and Qiyao Sun. 2025. [Rag-mcp: Mitigating prompt bloat in llm tool selection via retrieval-augmented generation](#). *arXiv preprint arXiv:2505.03275*.
- Luyu Gao, Xueguang Ma, Jimmy Lin, and Jamie Callan. 2022. [Precise zero-shot dense retrieval without relevance labels](#). *arXiv preprint arXiv:2212.10496*.
- Tianyu Gao, Gautier Izacard, Patrick Lewis, Sebastian Riedel, and Wen-tau Yih. 2023. Ragas: Automated evaluation of retrieval-augmented generation. *arXiv preprint arXiv:2309.15217*.
- Lovisa Hagström, Sara Vera Marjanović, Haeun Yu, Arnav Arora, Christina Lioma, Maria Maistro, Pepa Atanasova, and Isabelle Augenstein. 2025. [A reality check on context utilisation for retrieval-augmented generation](#). *Preprint*, arXiv:2412.17031.
- Xu Han, Zhiyuan Chen, Zhiyuan Liu, and Maosong Sun. 2023. A survey on large language models: Generalization, efficiency, and deployment. *arXiv preprint arXiv:2304.13712*.
- Zhenyu Ji, Nayeon Lee, Rudolf Frieske, and 1 others. 2023. [Survey of hallucination in natural language generation](#). *Transactions of the Association for Computational Linguistics*.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2019. [Billion-scale similarity search with GPUs](#). *IEEE Transactions on Big Data*. See also the FAISS arXiv preprint 2017.
- Omar Khattab and Matei Zaharia. 2020. [ColBERT: Efficient and effective passage search via contextualized late interaction over BERT](#). In *SIGIR*.
- Junkyum Kim and Divya Mahajan. 2025. [Vectorlitrage: Latency-aware and fine-grained resource partitioning for efficient rag](#). *Preprint*, arXiv:2504.08930.
- Youngwoo Kim, Razieh Rahimi, and James Allan. 2024. Discovering biases in information retrieval models using relevance thesaurus as global explanation. <https://arxiv.org/abs/2410.03584>. ArXiv preprint arXiv:2410.03584 [cs.IR].
- Abhishek Kumar, Ritesh Raj, and Manish Singh. 2024a. A comparative study of retrieval-augmented generation and fine-tuned transformer models for domain-specific chatbots. *Journal of Artificial Intelligence Research and Applications*.
- Abhishek Kumar, Manish Singh, and Ritesh Raj. 2024b. Rag vs. fine-tuning pipelines: Trade-offs and a case study on agricultural question answering. In *Proceedings of the 31st International Conference on Computational Linguistics (COLING)*.
- LangChain. 2024. State of ai agents 2024: A survey of over 1,300 professionals. <https://www.langchain.com/stateofaiagents>. Accessed: November 2025.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Chin-Yew Lin. 2004. [ROUGE: A package for automatic evaluation of summaries](#). In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.
- Qianren Mao, Qili Zhang, Hanwen Hao, Zhentao Han, Runhua Xu, Weifeng Jiang, Qi Hu, Zhijun Chen, Tyler Zhou, Bo Li, Yangqiu Song, Jin Dong, Jianxin Li, and Philip S. Yu. 2025. [Privacy-preserving federated embedding learning for localized retrieval-augmented generation](#). <https://arxiv.org/abs/2504.19101>. ArXiv preprint arXiv:2504.19101 [cs.CL].

- Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. [Factscore: Fine-grained atomic evaluation of factual precision in long form text generation](#). *Preprint*, arXiv:2305.14251.
- Anh Nguyen-Duc, Hoang Pham, Thuy Le, and 1 others. 2025. An empirical study of multi-agent rag for real-world university admissions counseling. In *arXiv preprint arXiv:2507.11272*.
- Rodrigo Nogueira and Kyunghyun Cho. 2019. [Pasage re-ranking with BERT](#). *arXiv preprint arXiv:1901.04085*.
- Jie Ouyang, Tingyue Pan, Mingyue Cheng, Ruiran Yan, Yucong Luo, Jiaying Lin, and Qi Liu. 2025. [Hoh: A dynamic benchmark for evaluating the impact of outdated information on retrieval-augmented generation](#). <https://arxiv.org/abs/2503.04800>. ArXiv:2503.04800 [cs.CL].
- Long Ouyang, Jeff Wu, Xu Jiang, and 1 others. 2022. [Training language models to follow instructions with human feedback](#). In *NeurIPS*.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.
- Rahul Patel, Lin Wang, and Abhinav Sharma. 2024. Financebench: Benchmarking rag for financial question answering. In *NeurIPS Workshop on Financial AI*.
- Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. [In-context retrieval-augmented language models](#). *Preprint*, arXiv:2302.00083.
- Nir Ratner, Yoav Levine, Yonatan Belinkov, Ori Ram, Inbal Magar, Omri Abend, Ehud Karpas, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. [Parallel context windows for large language models](#). *Preprint*, arXiv:2212.10947.
- Timo Schick, Felix Schütz, Pius von Dollen Seidl, and 1 others. 2023. [Toolformer: Language models can teach themselves to use tools](#). *arXiv preprint arXiv:2302.04761*.
- Sanat Sharma, DavidSeunghyun Yoon, Franck Dernoncourt, Dewang Sultania, Karishma Bagga, Mengjiao Zhang, Trung Bui, and Varun Kotte. 2024. [Retrieval augmented generation for domain-specific question answering](#). *arXiv preprint arXiv:2404.14760*. Workshop on Scientific Document Understanding (AAAI’24 workshop), February 20–27, 2024, Vancouver, Canada.
- Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. 2020. [Is multihop QA in DiRe condition? measuring and reducing disconnected reasoning](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8846–8863, Online. Association for Computational Linguistics.
- Sourav Verma. 2024. [Contextual compression in retrieval-augmented generation for large language models: A survey](#). *Preprint*, arXiv:2409.13385.
- Yuxuan Wang, Huy Nguyen, Nada Al-Khouri, and 1 others. 2024. A rag chatbot for precision medicine of multiple myeloma. *bioRxiv*.
- Lee Xiong, Chenyan Xiong, Ye Li, Kwok-Fung Tang, Jialin Liu, Paul Bennett, Junaid Ahmed, and Arnold Overwijk. 2020. [Approximate nearest neighbor negative contrastive learning for dense text retrieval](#). *Preprint*, arXiv:2007.00808.
- Jingbo Yang, Bairu Hou, Wei Wei, Yujia Bao, and Shiyu Chang. 2025. [Kvlink: Accelerating large language models via efficient kv cache reuse](#). *Preprint*, arXiv:2502.16002.
- Shunyu Yao, Jeffrey Yu, Dian Zhao, and 1 others. 2023. [ReAct: Synergizing reasoning and acting in language models](#). *arXiv preprint arXiv:2210.03629*.
- Hao Yu, Aoran Gan, Kai Zhang, Shiwei Tong, Qi Liu, and Zhaofeng Liu. 2025. [Evaluation of Retrieval-Augmented Generation: A Survey](#), page 102–120. Springer Nature Singapore.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [Bertscore: Evaluating text generation with bert](#). *Preprint*, arXiv:1904.09675.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging llm-as-a-judge with mt-bench and chatbot arena](#). *Preprint*, arXiv:2306.05685.
- Helena Zubkova, Ji-Hoon Park, and Sang-Woo Lee. 2025. [Sugar: Leveraging contextual confidence for smarter retrieval](#). *Preprint*, arXiv:2501.04899. ArXiv preprint arXiv:2501.04899.